CS480 – Introduction to Artificial Intelligence

Topic: Introduction
Chapter: 1

Mustafa Bilgic

http://www.cs.iit.edu/~mbilgic

https://twitter.com/bilgicm

AI is Everywhere Now

Email Filtering

"Wow! I've got one from some

Search Engines

Recommender Systems

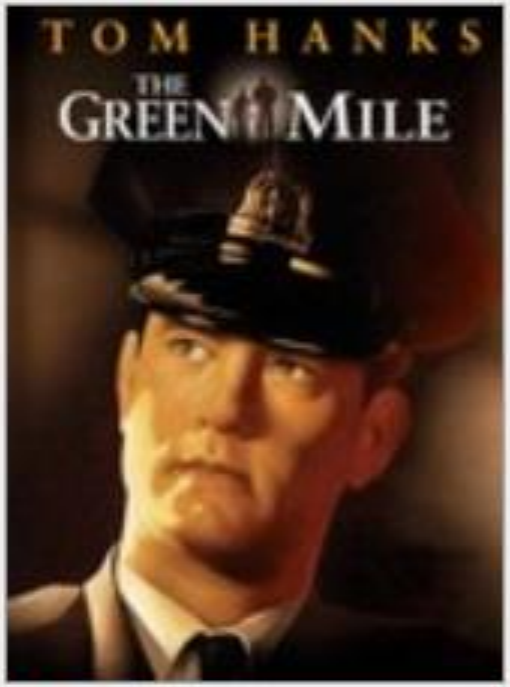

Add

Not Interested

The Green Mile

Because you enjoyed:

The Shawshank Redemption: Special Edition

Forrest Gump

Rain Man

Face Detection & Recognition

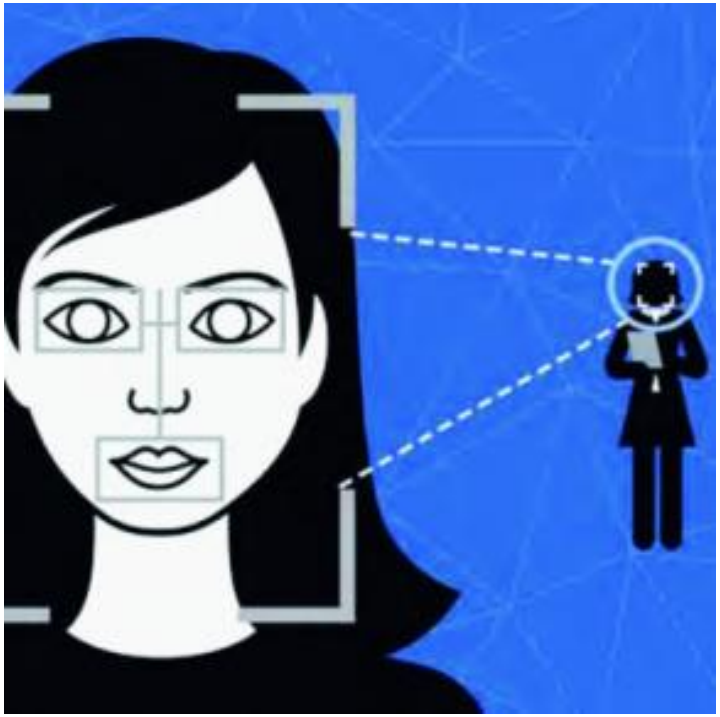

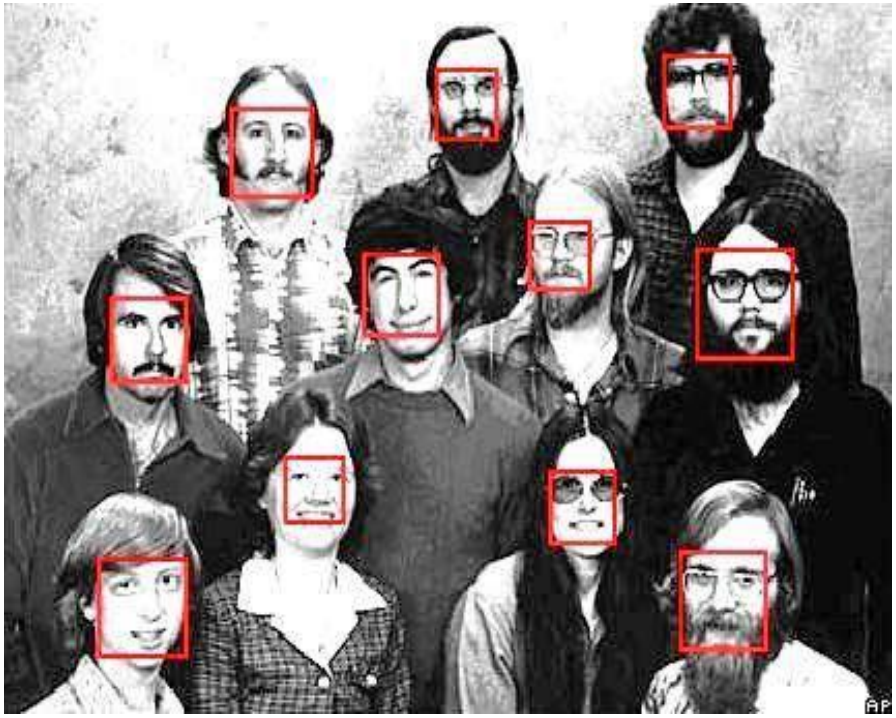

Medical Diagnosis

Intelligent Personal Assistants

Image Recognition + Translation

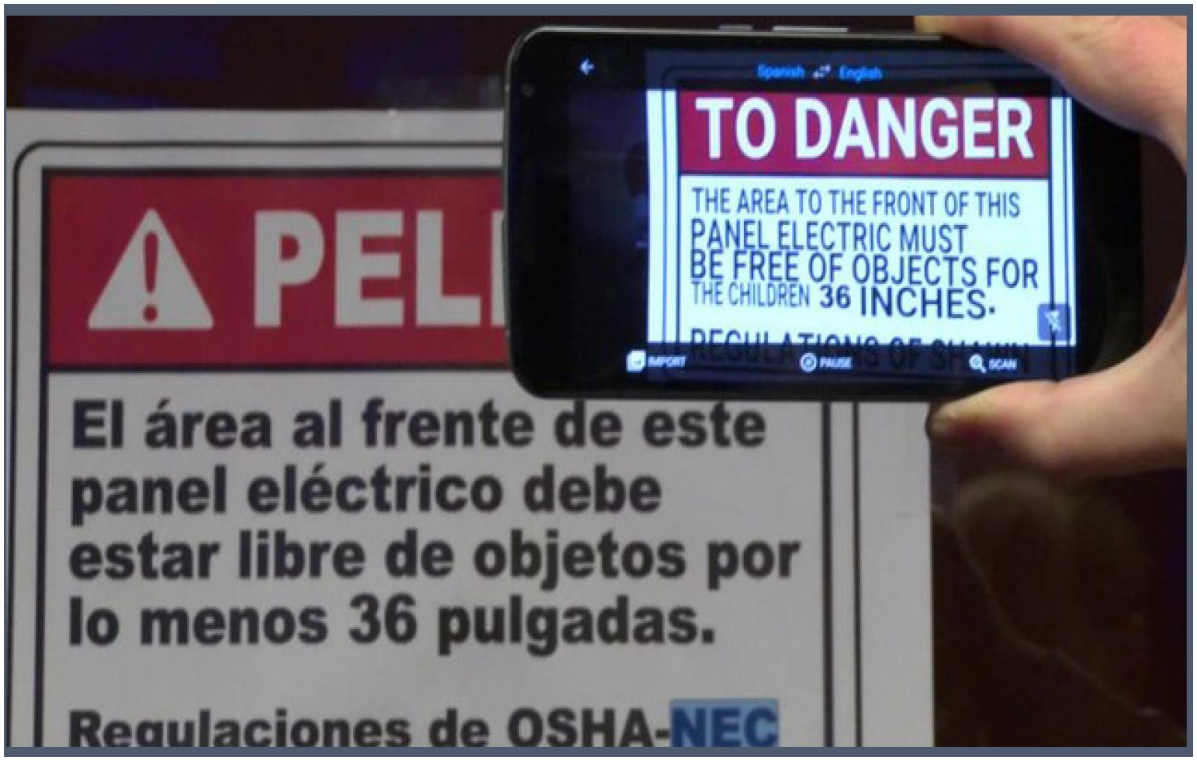

Self-driving Cars

… and of course Games!

Kasparov vs Deep Blue –1997

IBM Watson – Jeopardy – 2011

Google DeepMind – Go – 2016

CMU – Poker – 2017

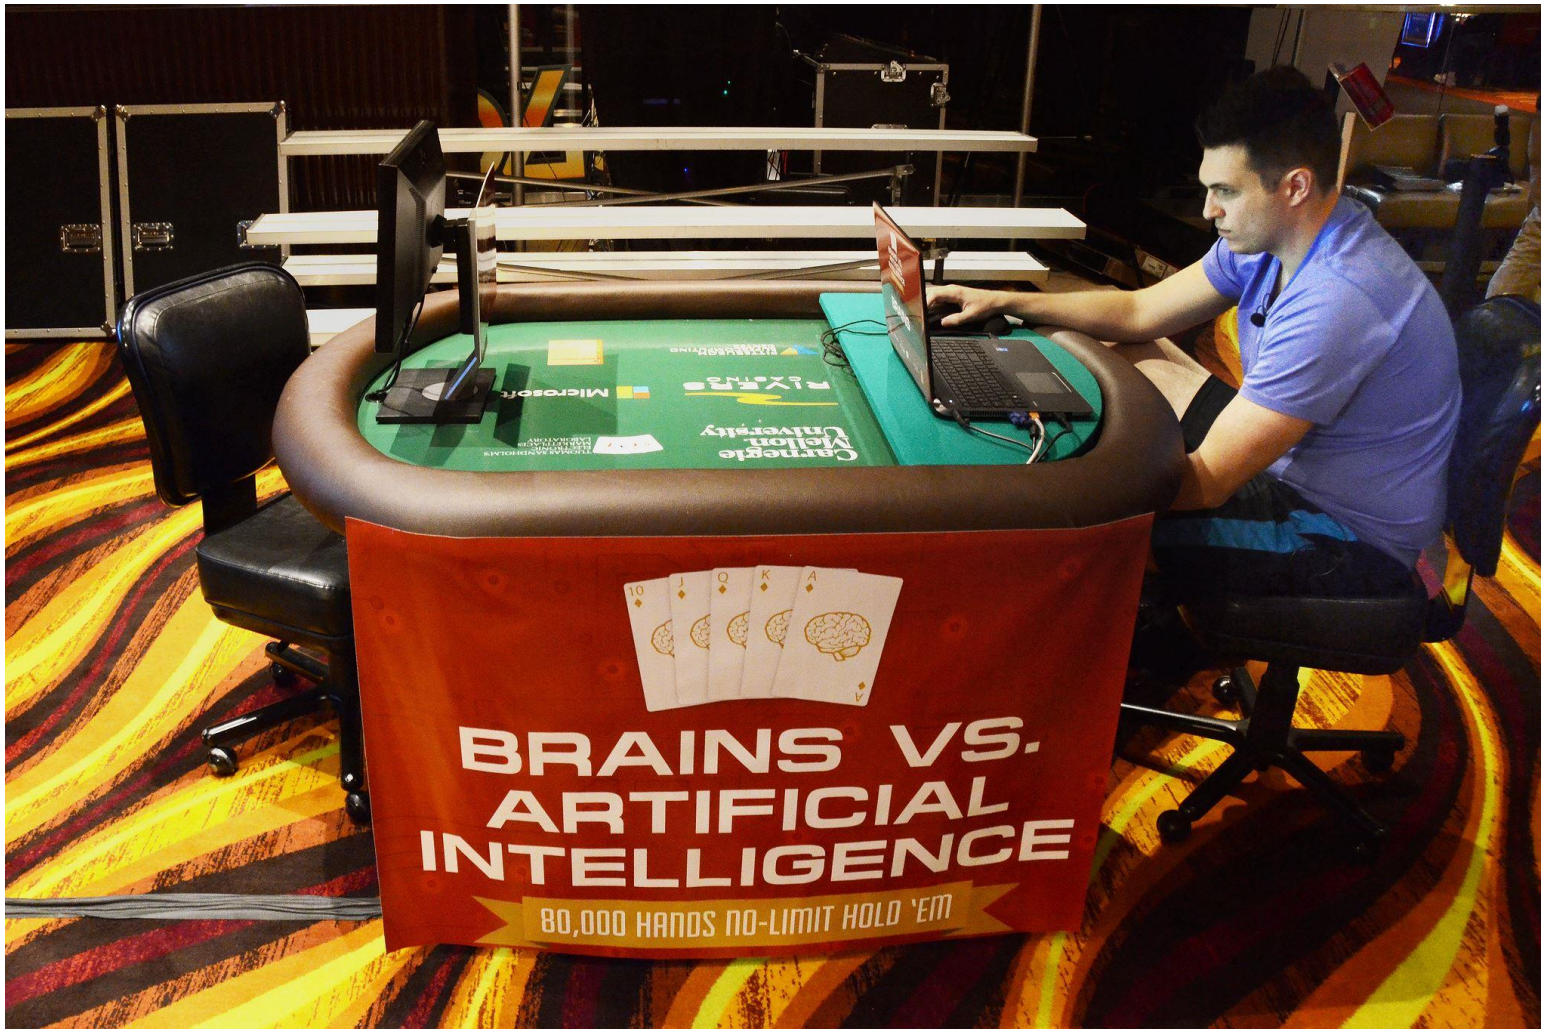

Any Other Examples

- ?

So, What is Artificial Intelligence?

- First, what is "intelligence"?
- Google
  - "The ability to acquire and apply knowledge and skills."
- Dictionary
  - "Capacity for learning, reasoning, understanding, and similar forms of mental activity; aptitude in grasping truths, relationships, facts, meanings, etc."

What is AI?

- https://www.lexico.com/en/definition/artificial_intelligence
  - "The theory and development of computer systems able to perform tasks normally requiring human intelligence, such as visual perception, speech recognition, decision-making, and translation between languages."
- https://www.merriam-webster.com/dictionary/artificial%20intelligence
  - "a branch of computer science dealing with the simulation of intelligent behavior in computers"
  - "the capability of a machine to imitate intelligent human behavior"
- https://www.britannica.com/technology/artificial-intelligence
  - "the ability of a digital computer or computer-controlled robot to perform tasks commonly associated with intelligent beings."
- https://en.wikipedia.org/wiki/Artificial_intelligence
  - "is intelligence demonstrated by machines, in contrast to the natural intelligence displayed by humans"

Intelligent?

- Calculators?

Intelligent?

- Search engines?

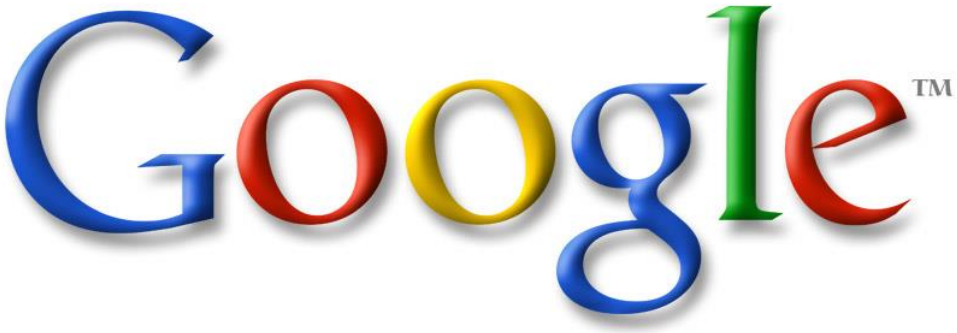

Imitate Humans?

- Would you call a robot that can perfectly imitate a human *intelligent*?

THE AI EFFECT

- “Every time we figure out a piece of it, it stops being magical; we say, 'Oh, that's just a computation.”
- “AI is whatever hasn't been done yet.”

Quotes from https://en.wikipedia.org/wiki/AI_effect

Humanly vs. rationally & thinking vs. acting

	Humanly	Rationally
Think	Thinking humanly	Thinking rationally
Act	Acting humanly	Acting rationally

Acting Humanly – the Turing test

- The imitation game
  - An operational test

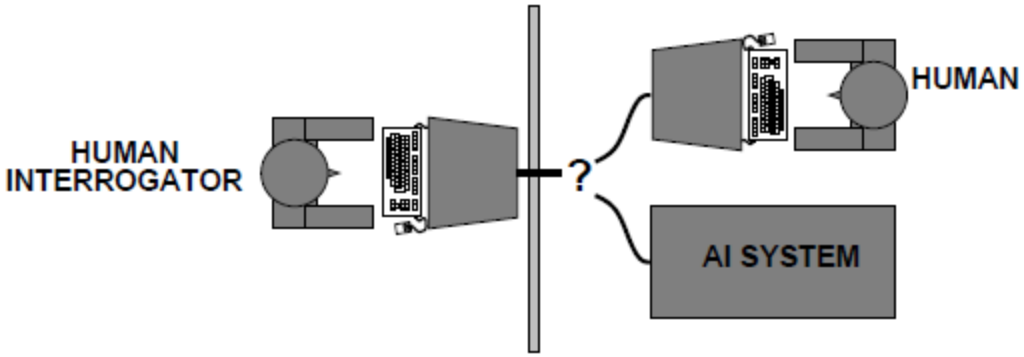

- The AI system needs to have:
  - Natural language processing
  - Knowledge representation
  - Automated reasoning
  - Machine learning

Thinking humanly – cognitive modeling

- Need to know how humans think
  - Introspection
  - Psychological experiments
  - Brain imaging
- Cognitive science
  - Based on experimental investigation of humans and animals

Thinking rationally – laws of thought

- Codify "right thinking"
  - Aristotle
- Logic
  - "Socrates is a man; all men are mortal; therefore, Socrates is mortal"
- Two main challenges
  - It is hard to encode esp. uncertain knowledge in formal logic
  - Can be computationally very demanding, unless it is provided some guidance

Acting rationally

- A **rational agent** is an agent that acts so as to achieve the best outcome, or when there is uncertainty, the best expected outcome.
- Two advantages
  - More general than thinking rationally, because acting rationally requires thinking rationally
  - More amenable to scientific development than the approaches based on human

Can machines think?

"The question of whether machines can think … is about as relevant as the question of whether submarines can swim."

Edsger Dijkstra (1984)

Intelligence and

- Consciousness
- Emotions
- Kindness
- Sense of humor
- Tell right from wrong
- Love
- Creativity
- Learning

A Great Read

- Turing, A. (1950). Computing machinery and intelligence. *Mind,* 59, 433-460.

Arguments against true AI & Counter Arguments

Source: Turing, A. (1950). Computing machinery and intelligence.

- The theological objection
  - Thinking is a function of man's immortal soul.
- The "Heads in the Sand" objection
  - The consequences of machines thinking would be too dreadful. Let us hope and believe they cannot do so.
- The mathematical objection
  - There are limitations to the power of discrete-state machines
- The argument from consciousness
  - "Not until a machine can write a sonnet or compose a concerto because of thoughts and emotions felt, and not by the chance of fall of symbols, could we agree that machine equals brain – that is, not only write it but know that it had written it."

Arguments against true AI & Counter Arguments

Source: Turing, A. (1950). Computing machinery and intelligence.

- Arguments about cannot do's
  - I grant you that you can make machines do all the things you have mentioned but you will never be able to make on do X.
    - Be kind, resourceful, beautiful, and friendly
    - Have initiative
    - Have a sense of humor
    - Fall in love
    - Enjoy strawberries and cream
    - Do something really new

Arguments against true AI

Source: You.

- ?

Weak vs Strong AI

- Weak AI
  - Build AI systems that are really good at one task
  - Most, if not all, of the current systems
- Strong AI
  - Build AI systems that are generally intelligent
  - Challenge: the whole is greater than the sum of its parts

The foundations - I

- Philosophy
  - Logic, induction, rationalism, empiricism
- Mathematics
  - Probability, statistics
- Computing
  - Algorithms, data
- Engineering
  - Chips, sensors, robotics

The foundations - II

- Economics
  - Utility, decision theory, game theory
- Neuroscience
  - The study of the brain
- Psychology
  - Behaviorism, cognitive psychology, how humans and animals think and act
- Linguistics
  - Grammar, syntax, how language relates to thinking

Subfields of AI

1. Communication and Perception
   - Language, speech, vision, robotics
2. Knowledge representation and reasoning
   - Logic, probability, planning, decision making
3. Learning
   - Machine learning
4. Problem solving
   - Search, constraint satisfaction, game playing

AI vs ML vs DL

- A common misconception
  - AI = Machine Learning = Deep Learning
- Reality
  - Deep Learning $\subset$ Machine Learning $\subset$ AI

Machine Learning

Developing programs that improve their performance through experience at a given task

Tom Mitchell, Machine Learning

A few ML Examples

- Face recognition
- Speech recognition
- Game playing
- Medical diagnosis
- Scientific data analysis
- Behavior analysis
- Product recommendations
- Ad placements
- Personalization
- Credit scoring
- Fraud detection
- …

History - I

- Gestation:1943 – 1955
  - Based on:
    - Physiology and function of the neurons in the brain
    - Formal analysis of propositional logic
    - Theory of computation
  - First neural network computer – 1950
  - Turing test – 1950
- Birth: 1956
  - Darthmouth workshop: the term AI was coined
  - Logic Theorist – was able to prove most theorems in the book *Principia Mathematica*

History - II

- Early enthusiasm: 1952 – 1969
  - General Problem Solver (GPS) – imitate human problem-solving protocols – thinking humanly approach
  - Geometry Theorem Prover – was able to prove theorems that many math students found to be tricky
  - Checkers – the program that learned to play checkers – disproved the idea that the computers can do only what they are told to do
  - Lisp – the dominant AI programming language for about 30 years
  - Many microworlds – limited domains
    - SAINT – solved closed-form calculus integration problems
    - ANALOGY – solved geometric analogy problems that appear in IQ tests
    - STUDENT – solved algebra story problems
  - Perceptrons

History - III

- A dose of reality: 1966 – 1973
  - Merely syntactic manipulations
    - "The spirit is willing but the flesh is weak" => "The vodka is good but the meat is rotten"
  - Intractability
    - Tried many possible combinations till worked
    - Worked initially because microworlds contained very few objects and actions
  - Representation limitations of perceptrons
    - Almost killed the neural net research until 80s

History - IV

- Knowledge-based systems: 1969 – 1979
  - Narrow areas of expertise with domain knowledge integration
  - DENDRAL – inferred molecular structure
    - Integrated domain knowledge to guide and limit the search
  - MYCIN – diagnosed blood infections
    - Was better than junior doctors
    - Was able to handle uncertain knowledge
  - Developments in knowledge representation

History - V

- AI becomes industry: 1983 – present
  - R1 – the first successful commercial application – 1982
    - Helped configure orders for new computer systems
    - By 1986, it saved an estimated of $40 million a year
  - AI industry
    - 1980 – a few million dollars
    - 1988 – billions of dollars

History - VI

- Return of neural networks: 1986 – present
- AI adopts scientific method: 1987 – present
- Emergence of intelligent agents: 1995 – present
- Very large datasets: 2001 – present

The book does not have the most recent history, for obvious reasons. What do you think the recent history of AI should include?

The state of the art w.r.t. the book

- Robotic vehicles
  - Driverless cars
- Speech recognition
  - Speech-to-text
  - Automated dialog management systems
- Autonomous planning and scheduling
  - NASA's mission robots
- Game playing
  - Deep blue
- Spam filtering
- Machine translation

The state of the art – Others

- Personal agents
  - E.g., Siri, Google now, Cortana, Alexa
- Recommendations
  - E.g., Netflix, Amazon
- Face detection and recognition
  - Even personal cameras can do face detection now
- Near real-time speech recognition & translation
  - E.g., Skype translator
- Games
  - E.g., IBM Watson – Jeopardy, AlphaGo – Go , CMU – Poker

The state of the art

- What do you think some of the recent AI accomplishments and cool applications are

AI Winter(s)?

- 1966
  - National Research Council report: “machine translation was more expensive, less accurate and slower than human translation”
- 1969
  - “Perceptrons” book; showed the limits of perceptrons, the building blocks of neural networks
- 1970s
  - The Lighthill report at UK; the problem of combinatorial explosion and intractability
  - Amendment to DARPA’s funding; required “mission-oriented” research rather than “basic” research
- 1987
  - The beginning of the collapse of the LIPS machine and expert systems
- 2020?

https://en.wikipedia.org/wiki/AI_winter

What is new?

1. Data
   - We generate **so** much data
   - We can and do store **all** of it
2. Computing power
   - Moore's law: "the number of transistors in a integrated dense circuit doubles about every two years"
   - GPU computation

DATA NEVER SLEEPS 6.0

How much data is generated *every minute?*

There's no way around it: big data just keeps getting bigger. The numbers are staggering, but they're not slowing down. By 2020, it's estimated that for every person on earth, 1.7 MB of data will be created every second. In our 6th edition of Data Never Sleeps, we once again take a look at how much data is being created all around us every single minute of the day—and we have a feeling things are just getting started.

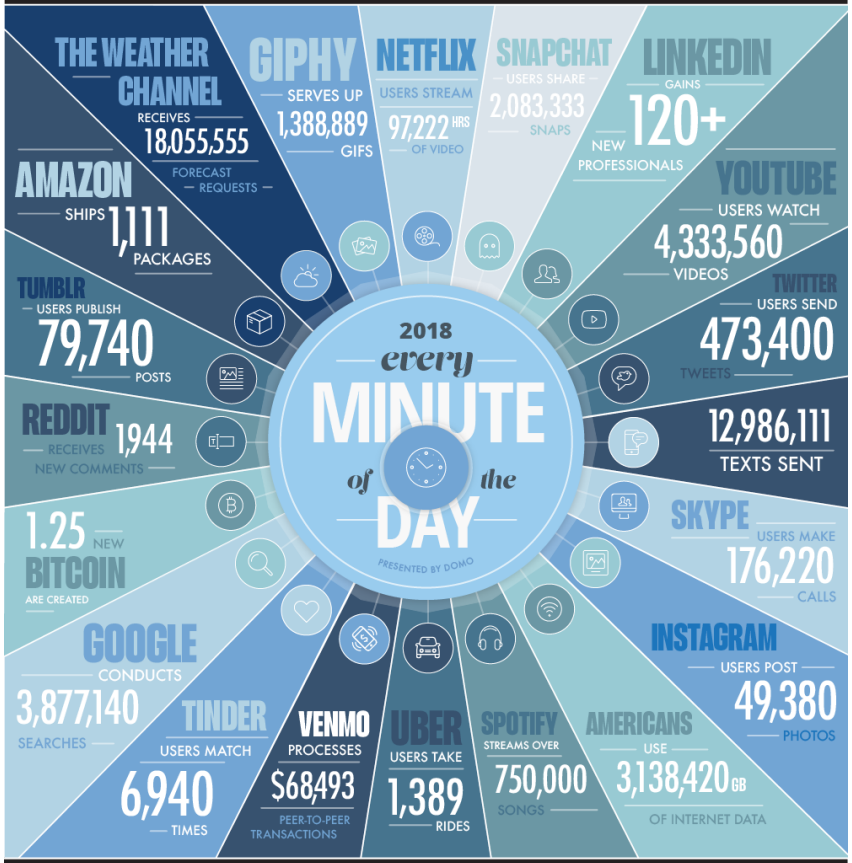

The world's internet population is growing significantly year-over-year. In 2017, internet usage reached 47% of the worlds population and now represents 3.8 billion people.

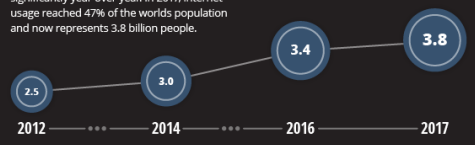

GLOBAL INTERNET POPULATION GROWTH 2012–2017
(IN BILLIONS)

The ability to make data-driven decisions is crucial to any business. With each click, swipe, share, and like, a world of valuable information is created. Domo puts the power to make those decisions right into the palm of your hand by connecting your data and your people at any moment, on any device, so they can make the kind of decisions that make an impact.

Learn more at domo.com

SOURCES: STATISTA, LINKEDIN, INTERNET LIVE STATS, EXPANDED RAMBLINGS, SLASH FILM, RIAA, BUSINESS OF APPS, INTERNATIONAL TELECOMMUNICATIONS UNION, INTERNATIONAL DATA CORPORATION

Moore's Law – The number of transistors on integrated circuit chips (1971-2016)

Our World in Data

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are strongly linked to Moore's law.

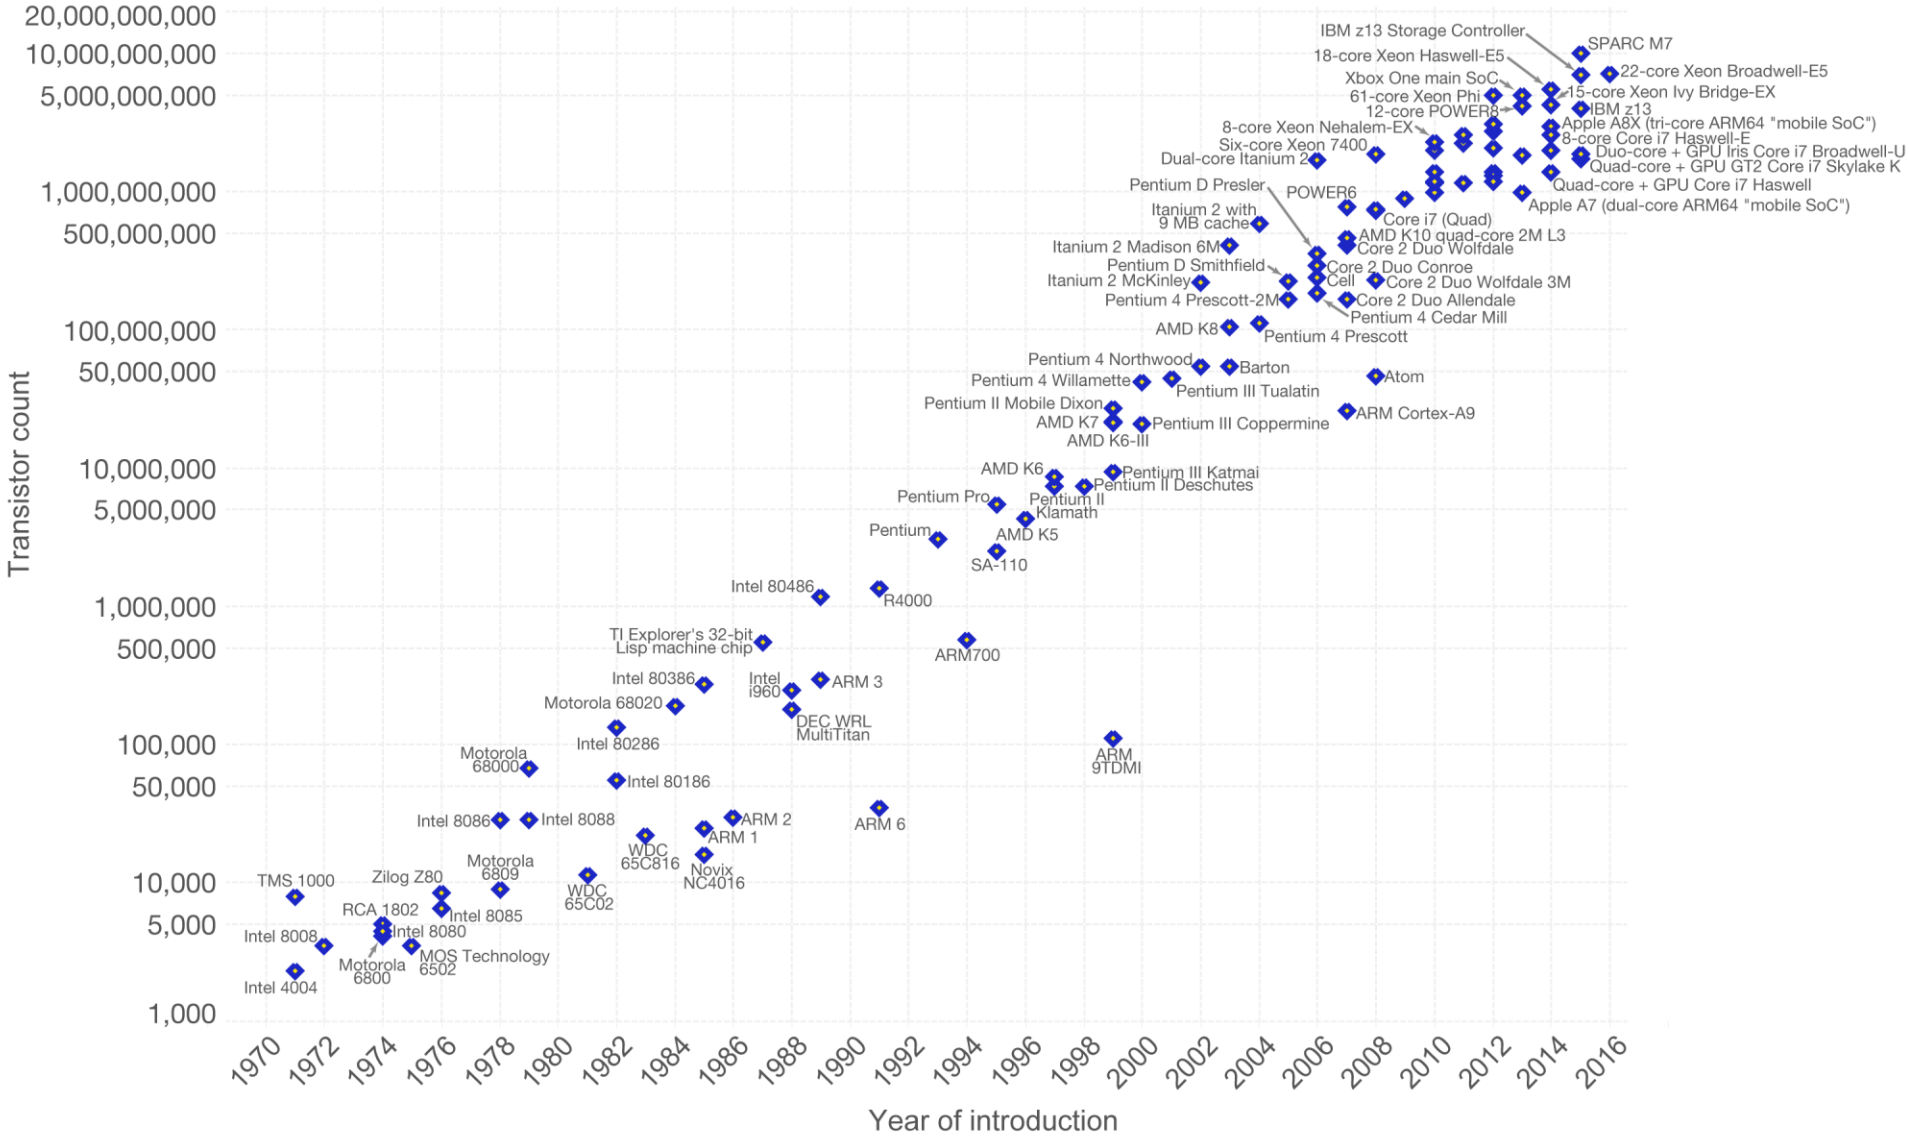

Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)

The data visualization is available at OurWorldinData.org. There you find more visualizations and research on this topic.

Some of the Current Challenges

- Privacy
- Fairness
- Accountability
- Transparency
- Misinformation
- Job loss and creation
- Regulations

Rest of the semester - I

- Intelligent agents – Chapter 2
  - Environment, performance, agent programs
- Search – Chapter 3, 5, & 6
  - Problem solving through uninformed and informed search
  - Game playing
  - Constraint satisfaction

Rest of the semester - II

- Knowledge representation and reasoning – Chapters 7, 8, & 9
  - Propositional logic
  - First-order logic
  - Resolution algorithm

Rest of the semester - III

- Uncertainty and Probabilistic Reasoning – Chapters 13 & 14
  - Probability theory
  - Bayesian networks
- Decision making – Chapter 16
  - Utility theory
  - Value of information

Rest of the semester - IV

- Learning – Chapters 18 & 20
  - Supervised learning
  - Decision trees
  - Naïve Bayes
  - Logistic regression
  - Neural networks