CS480 – Introduction to Artificial Intelligence

Topic: Search
Chapter: 3

Mustafa Bilgic

🔗 http://www.cs.iit.edu/~mbilgic

🐦 https://twitter.com/bilgicm

MOTIVATION

- We have discussed several types of agents
  - Reflex agents, goal-based agents, utility-based agents, etc.
- Reflex agents act purely based on current percept, which is limited
- We now discuss, in detail, the goal-based agents, where a sequence of actions are necessary to reach a desired world state
- Actions have associated costs with them and we'd like to minimize the total cost to reach the goal state

REPRESENTATION

- The world is represented as an atomic state

- An action takes us from one world state to another

- The environment is
  - Observable
  - Discrete
  - Deterministic
  - Known

An Example

- We are now in Chicago

- We'd like to drive to Pittsburgh

- We represent the world as "In X" where X is a member of a predefined set of major cities in US

- An action takes us from being in one major city (a world state) to being in a neighboring major city (another world state)

- The cost of the action is the distance traveled

- We would like to travel from Chicago to Pittsburgh while minimizing the distance traveled

THE SEARCH PROCESS

- We will start at an initial state

- At each state, there is a set of actions we can perform
  - The action that can be performed can be different depending on the current world state

- We have a test that can specify whether a given state is a goal state

- Each action has a cost

- We want to find the optimal solution, i.e., the minimum-cost sequence to reach the goal state

MORE FORMALLY

- The **initial state**
  - *In(Chicago)*

- A set of **actions**
  - *Go(Milwaukee), Go(Madison), Go(Bloomington)*, etc.

- A **transition model**: *Result(s, a): S→S*
  - *Result(In(Chicago), Go(Madison)) = In(Madison)*

- The **goal test**
  - *{In(Pittsburgh)}*

- A **path cost**: Sum of step costs *c(s,a,s')*
  - *c(In(Chicago), Go(Madison), In(Madison)) = 150*

ASSUMPTIONS AND SIMPLIFICATIONS

- For the Chicago to Pittsburgh problem, we made several assumptions and simplifications

- What are some of these?

The Vacuum World

- **States:** The cross-product of the agent's location and dirt locations
  - How many possible states?
- **Initial state:** Any state
- **Actions:** Left, Right, and Suck
- **Transition model:** Expected effects. Except Left on the left square, Right on the right square, and Suck on a clean square has no effect
- **Goal test:** Test whether all squares are clean
- **Path cost:** Each step costs 1; the path cost is then the total number of steps

The Vacuum State Space

The 8-Puzzle

7	2	4
5		6
8	3	1

Start State

|	1	2
3	4	5
6	7	8

Goal State

The 8-Puzzle

- **States:** All possible locations of the eight tiles and the blank tile
  - How many?
- **Initial state:** Any state
- **Actions:** Left, Right, Up, Down, except on the edges
- **Transition model:** The state with the blank the numbered tile switched
- **Goal test:** Blank on the top left, others are ordered in increasing order
- **Path cost:** Each step costs 1; the path cost is then the total number of steps

The 8-Queens Problem

THE 8-QUEENS PROBLEM

- **States:** Any arrangement of 0 to 8 queens on the board
  - How many?
- **Initial state:** No queens on the board
- **Actions:** Add a queen to an empty square
- **Transition model:** The board with the added queen
- **Goal test:** 8 queens on the board, none attacked
- **Path cost:** Each step costs 1; the path cost is then the total number of steps

ROMANIA

ROMANIA

- **States:** Any city

- **Initial state:** Arad

- **Actions:** Go to a neighboring city

- **Transition model:** The city that was just traveled

- **Goal test:** Bucharest

- **Path cost:** Step costs are distances between the neighboring cities; path cost is the total distance traveled

LET'S SOLVE TOY VERSIONS

○ The 4-Queens problem

- Initial state: Empty board

- Action: Place a queen on the left most empty column

○ The 3-Puzzle problem

- Initial state: First row: 2 – Blank, Second row: 3 – 1

- Action: Move the blank tile

SEARCH

- The key algorithm of the goal-based agent is *search*

1. Start with the initial state
2. Expand that state through possible actions
3. Pick an unexplored state
4. If goal, return solution; otherwise, go to step 2

TREE SEARCH

function TREE-SEARCH(*problem*) **returns** a solution, or failure
    initialize the frontier using the initial state of *problem*
    loop do
        if the frontier is empty **then return** failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state **then return** the corresponding solution
        expand the chosen node, adding the resulting nodes to the frontier

The key question is which leaf node to pick for expansion; in other words, what data structure is used for "frontier?" A FIFO queue? A LIFO queue? A priority queue? An unordered set? Are duplicates allowed?

Two Famous and Important Search Strategies

- Depth-first search
  - i.e., go as deep as possible
  - **frontier** = LIFO queue (stack)

- Breadth-first search
  - i.e., go broad; expand all the nodes at level i before expanding the nodes at level $i+1$
  - **frontier** = FIFO queue

DEPTH-FIRST SEARCH

BREADTH-FIRST SEARCH

Let's See Them in Action

- 4-Queens

- 3-Puzzle

- Romania

COMPLEXITY

- We're concerned with

  - Time complexity

  - Space complexity

  - Completeness (if there is a solution, can it find it?)

  - Optimality (is the found solution the optimal one?)

- Three relevant quantities

  - The number of nodes expanded

  - The maximum size of the **frontier**

  - The cost of the solution

ANALYSIS – TREE SEARCH

| Criterion	Depth-first	Breadth-first
Complete?		
Time		
Space		
Optimal?		

MEMORY IS A BIG PROBLEM

- Breadth-first search has an exponential memory requirement

- Depth-first search requires polynomial size memory, but it is not complete and it is not optimal

- What can we do?
  - Depth-limited search
    - Impose a limit on the depth
    - Is it complete? Optimal?
  - Iterative-deepening search
    - Start with a depth limit of 0, and run depth-limited search; if no goal found, increase the depth limit and repeat the process
    - Is it complete? Optimal?

ITERATIVE-DEEPENING SEARCH

# WHAT IF THE STEP COSTS ARE NOT IDENTICAL?

- BFS and Iterative Deepening are not guaranteed to be optimal anymore
- Define $f(n) = h(n) + g(n)$
  - $h(n)$: Cost estimate from n to the goal
  - $g(n)$: Cost from the initial state to n

1. Uniform-cost search
   - Expand using $g(n)$

2. Greedy best-first search
   - Expand using $h(n)$

3. A* search
   - Expand using $f(n)$

THE RUNNING EXAMPLE

Edge costs

h(n)

|			
Arad	366	**Mehadia**	241
Bucharest	0	**Neamt**	234
Craiova	160	**Oradea**	380
Drobeta	242	**Pitesti**	100
Eforie	161	**Rimnicu Vilcea**	193
Fagaras	176	**Sibiu**	253
Giurgiu	77	**Timisoara**	329
Hirsova	151	**Urziceni**	80
Iasi	226	**Vaslui**	199
Lugoj	244	**Zerind**	374

Uniform-Cost Search

- Arad to Bucharest

GREEDY BEST-FIRST SEARCH

(a) The initial state

Arad
366

(b) After expanding Arad

Arad

Sibiu   Timisoara   Zerind
253   329   374

(c) After expanding Sibiu

Arad

Sibiu   Timisoara   Zerind
               329   374

Arad   Fagaras   Oradea   Rimnicu Vilcea
366   176   380   193

(d) After expanding Fagaras

Arad

Sibiu   Timisoara   Zerind
               329   374

Arad   Fagaras   Oradea   Rimnicu Vilcea
366   380   193

Sibiu   Bucharest
253   0

A* Search

(a) The initial state

(b) After expanding Arad

(c) After expanding Sibiu

(d) After expanding Rimnicu Vilcea

(e) After expanding Fagaras

(f) After expanding Pitesti

GOAL TEST

Test when a goal state is generated or chosen for expansion?

Repeated States

- Repeated states can be a major problem

- Repeated states are especially common for problems where actions are reversible

  - But, they also occur for problems where there are multiple distinct paths between two states

- One way to avoid this problem is to remember the set of explored and generated states

Tree Search vs. Graph Search

function Tree-Search(*problem*) **returns** a solution, or failure
   initialize the frontier using the initial state of *problem*
   loop do
      if the frontier is empty **then return** failure
      choose a leaf node and remove it from the frontier
      if the node contains a goal state **then return** the corresponding solution
      expand the chosen node, adding the resulting nodes to the frontier

function Graph-Search(*problem*) **returns** a solution, or failure
   initialize the frontier using the initial state of *problem*
   initialize the explored set to be empty
   loop do
      if the frontier is empty **then return** failure
      choose a leaf node and remove it from the frontier
      if the node contains a goal state **then return** the corresponding solution
      add the node to the explored set
      expand the chosen node, adding the resulting nodes to the frontier
      only if not in the frontier or explored set

Warning: this is Figure 3.7 in the book. The last line needs to change to guarantee optimality.

MODIFIED GRAPH SEARCH

function GRAPH-SEARCH(*problem*) **returns** a solution, or failure
    initialize the frontier using the initial state of *problem*
    initialize the explored set to be empty
    loop do
        if the frontier is empty **then return** failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state **then return** the corresponding solution
        add the node to the explored set
        expand the chosen node
            for child in children of the chosen node:
                if child is not in explored:
                    add child to frontier only if child is not in frontier or
                    child has a lower cost than the one in frontier

GRAPH SEARCH EXAMPLES

QUESTIONS TO THINK ABOUT

How does graph search affect

- Completeness,

- Time complexity,

- Space complexity, and

- Optimality

of various search algorithms?

A* Optimal?

- Tree version is optimal if
  - *h(n)* is <u>admissible</u>
- Graph version is optimal if
  - *h(n)* is <u>consistent</u>

ADMISSIBLE? CONSISTENT?

- **Admissible** if
  - *h(n)* never overestimates the optimal cost
- **Consistent** if
  - If $h(n) \leq c(n,n') + h(n')$

A* Optimality – Graph Search – Proof

1. If $h(n)$ is consistent, then $f(n)$ along any path is non-decreasing

2. When n is expanded, the optimal path to it has been found

- Proof is given on page 95

A* vs uniform-cost search

- Compare and contrast A* and UCS
  - Which one is better and why?
- Given the same heuristic function, can you design a complete and optimal algorithm that expands fewer nodes than A*?

8-Puzzle Heuristics

- h_1 = The number of misplaced tiles

- h_2 = Manhattan distance

Let's see an example

5	6	1
2	4	8
3	7	

EFFECTIVE BRANCHING FACTOR

- If the total number of nodes generated is N and the solution depth is d, then
  - $b*$ is the branching factor that a uniform tree of depth d would need to have in order to contain $N+1$ nodes
- $N+1 = 1 + b* + (b*)^2 + \ldots + (b*)^d$
- If A* finds a solution at depth 4 using 40 nodes, what is $b*$?
  - ≈2.182
- A good heuristic function achieves $b*≈1$

IDS vs A*

	Search Cost (nodes generated)			Effective Branching Factor		
d	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Figure 3.29 Comparison of the search costs and effective branching factors for the ITERATIVE-DEEPENING-SEARCH and A* algorithms with h_1, h_2. Data are averaged over 100 instances of the 8-puzzle for each of various solution lengths d.

HEURISTIC FUNCTIONS

- Let $h(n)=0$ for all n
  - Is it admissible?
  - Is it consistent?
  - Is it any good?

- Can we say an admissible heuristic function h_i is *always* better than another admissible heuristic function h_j?

- How can we find good heuristics?

- A heuristic: run uniform-cost search, find the solution, return the cost of the solution as the heuristic value. Can you beat it?

- Let $h(n) = \max(h_1(n), h_2(n), \ldots, h_m(n))$, where $h_i(n)$, are admissible
  - Is $h(n)$ admissible?
  - Can we say $h(n)$ is better than any of $h_i(n)$?
  - What's the catch?

Summary

- Depth-first search

- Breadth-first search

- Depth-limited depth-first search

- Iterative deepening

- Uniform cost search: $g(n)$

- Greedy best-first search: $h(n)$

- A* search: $g(n) + h(n)$

  Which one is better under what conditions?

NEXT

- Skip Chapter 4

- Chapter 5: Adversarial search – game playing