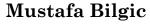
CS480 – Introduction to Artificial Intelligence

TOPIC: SEARCH

CHAPTER: 3





http://www.cs.iit.edu/~mbilgic



https://twitter.com/bilgicm

OUTLINE

- Problem solving agent
- The search process
- Tree search
- Graph search
- Complexity analysis

MOTIVATION

- We have discussed several types of agents
 - Reflex agents, goal-based agents, utility-based agents, etc.
- Reflex agents act purely based on current percept, which is limited
- We now discuss, in detail, the goal-based agents, where a sequence of actions are necessary to reach a desired world state
- Actions have associated costs with them and we'd like to minimize the total cost to reach the goal state

REPRESENTATION

- The world is represented as an atomic state
- An action takes us from one world state to another
- The environment is
 - Observable
 - Discrete
 - Deterministic
 - Known

AN EXAMPLE

- We are now in Chicago
- We'd like to drive to Pittsburgh
- We represent the world as "In X" where X is a member of a predefined set of major cities in US
- An action takes us from being in one major city (a world state) to being in a neighboring major city (another world state)
- The cost of the action is the distance traveled
- We would like to travel from Chicago to Pittsburgh while minimizing the distance traveled

THE SEARCH PROCESS

- We will start at an initial state
- At each state, there is a set of actions we can perform
 - The action that can be performed can be different depending on the current world state
- We have a test that can specify whether a given state is a goal state
- Each action has a cost
- We want to find the optimal solution, i.e., the minimum-cost sequence to reach the goal state

More Formally

- The initial state
 - In(Chicago)
- A set of actions
 - Go(Milwaukee), Go(Madison), Go(Bloomington), etc.
- A transition model: $Result(s, a): S \rightarrow S$
 - Result(In(Chicago), Go(Madison)) = In(Madison)
- The goal test
 - $\{In(Pittsburgh)\}$
- A path cost: Sum of step costs c(s,a,s')
 - c(In(Chicago), Go(Madison), In(Madison)) = 150

ASSUMPTIONS AND SIMPLIFICATIONS

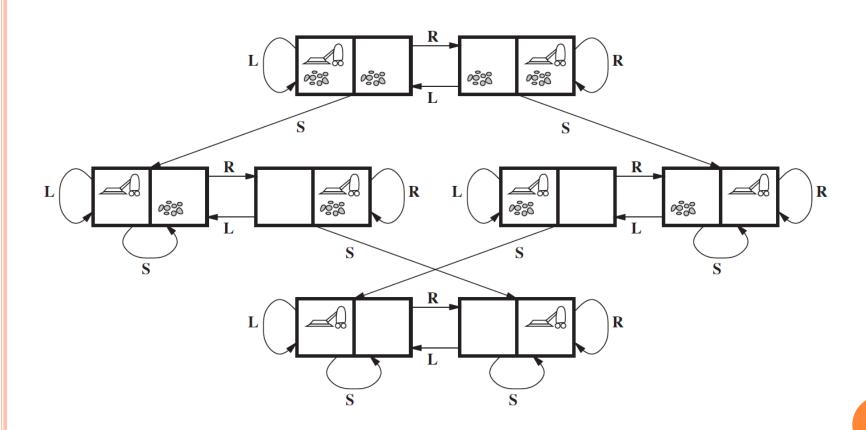
- For the Chicago to Pittsburgh problem, we made several assumptions and simplifications
- o What are some of these?

THE VACUUM WORLD

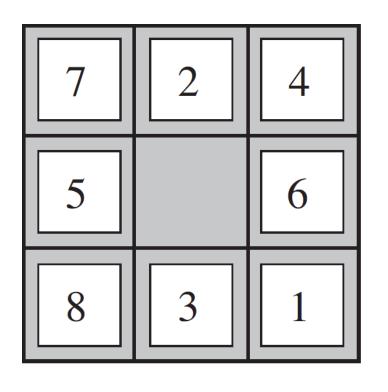
- **States:** The cross-product of the agent's location and dirt locations
 - How many possible states?
- Initial state: Any state
- Actions: Left, Right, and Suck
- Transition model: Expected effects. Except Left on the left square, Right on the right square, and Suck on a clean square has no effect
- Goal test: Test whether all squares are clean
- Path cost: Each step costs 1; the path cost is then the total number of steps



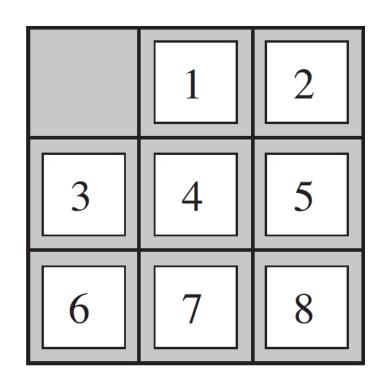
THE VACUUM STATE SPACE



THE 8-PUZZLE



Start State



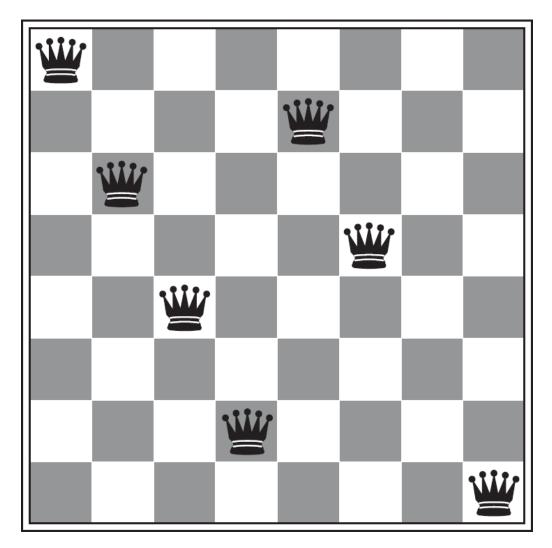
Goal State

THE 8-PUZZLE

- States: All possible locations of the eight tiles and the blank tile
 - How many?
- Initial state: Any state
- Actions: Left, Right, Up, Down, except on the edges
- Transition model: The state with the blank the numbered tile switched
- Goal test: Blank on the top left, others are ordered in increasing order
- Path cost: Each step costs 1; the path cost is then the total number of steps



The 8-Queens Problem

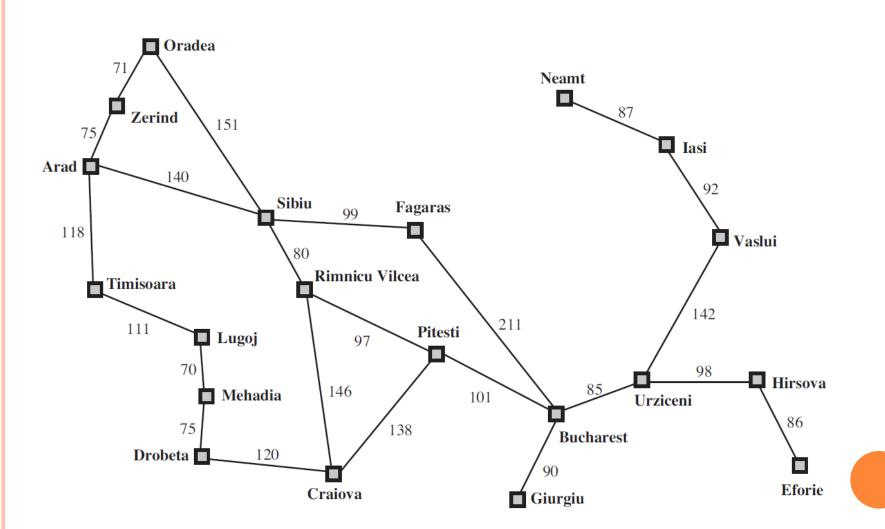


THE 8-QUEENS PROBLEM

- States: Any arrangement of 0 to 8 queens on the board
 - How many?
- Initial state: No queens on the board
- Actions: Add a queen to an empty square
- Transition model: The board with the added queen
- Goal test: 8 queens on the board, none attacked
- Path cost: Each step costs 1; the path cost is then the total number of steps



ROMANIA



ROMANIA

- States: Any city
- o Initial state: Arad
- Actions: Go to a neighboring city
- Transition model: The city that was just traveled
- o Goal test: Bucharest
- Path cost: Step costs are distances between the neighboring cities; path cost is the total distance traveled

LET'S SOLVE TOY VERSIONS

- The 4-Queens problem
 - Initial state: Empty board
 - Action: Place a queen on the left most empty column
- The 3-Puzzle problem
 - Initial state: First row: 2 Blank, Second row: 3 1
 - Action: Move the blank tile

SEARCH

- The key algorithm of the goal-based agent is search
- 1. Start with the initial state
- 2. Expand that state through possible actions
- 3. Pick an unexplored state
- 4. If goal, return solution; otherwise, go to step 2

TREE SEARCH

function TREE-SEARCH(problem) returns a solution, or failure initialize the frontier using the initial state of problem loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier if the node contains a goal state then return the corresponding solution expand the chosen node, adding the resulting nodes to the frontier

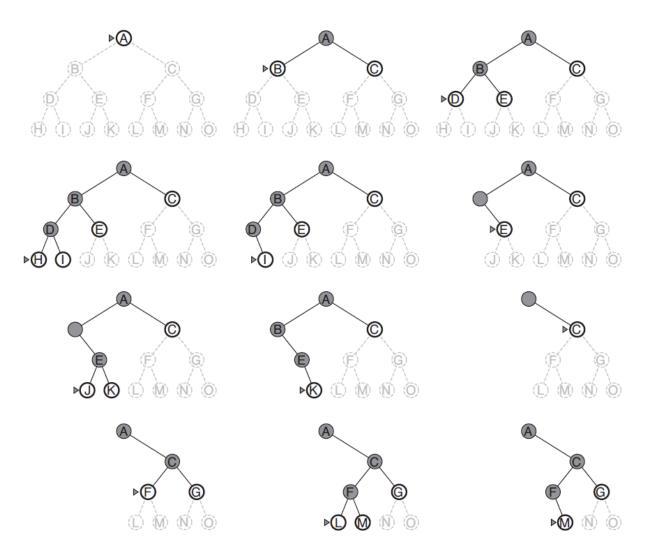
The key question is which leaf node to pick for expansion.

TWO FAMOUS AND IMPORTANT SEARCH STRATEGIES

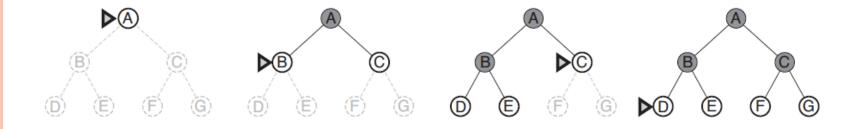
- Depth-first search
 - i.e., go as deep as possible
 - frontier = LIFO queue (stack)
- Breadth-first search
 - i.e., go broad; expand all the nodes at level i before expanding the nodes at level i+1
 - **frontier** = FIFO queue

This is the textbook version; we do it a bit differently. Why?

DEPTH-FIRST SEARCH

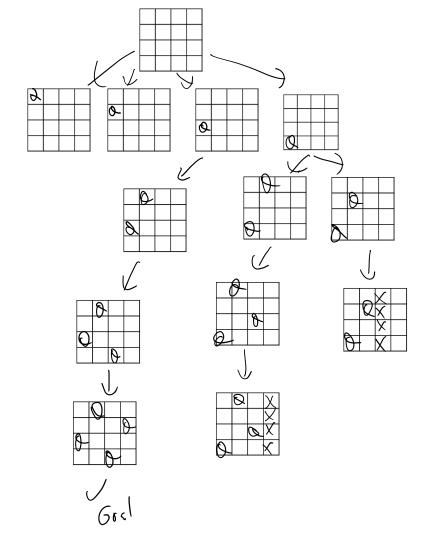


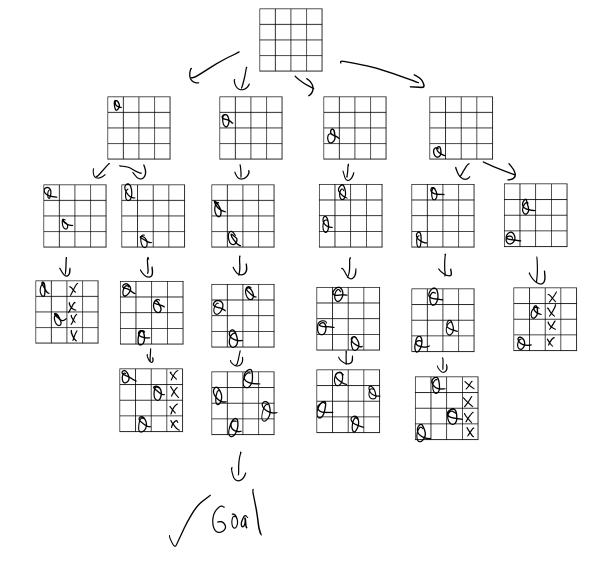
Breadth-first Search

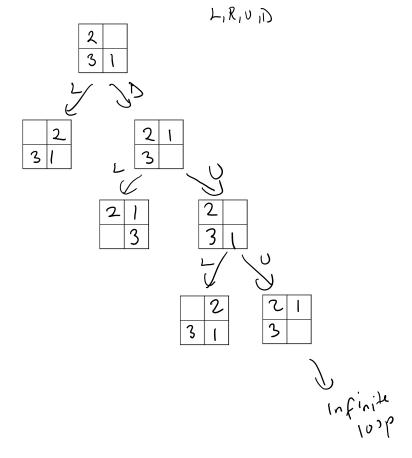


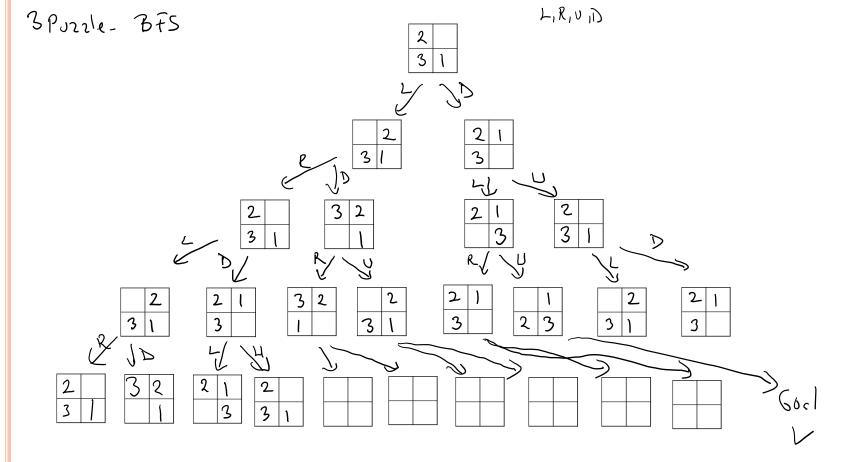
LET'S SEE THEM IN ACTION

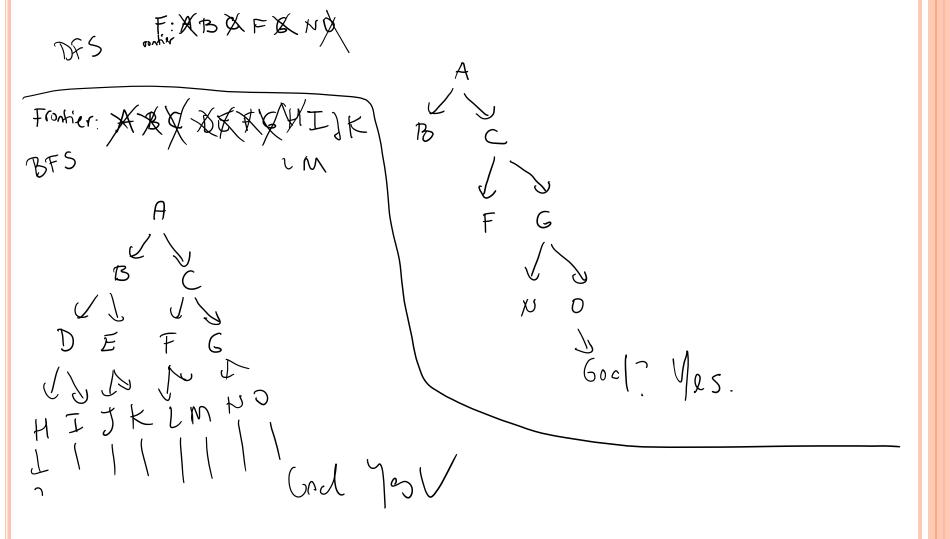
- 4-Queens
- 3-Puzzle
- Travel
 - $A \rightarrow B, C$
 - $B \rightarrow D, E$
 - $C \rightarrow F, G$
 - $D \rightarrow H, I$
 - $E \rightarrow J, K$
 - $F \rightarrow L, M$
 - $G \rightarrow N, O$
 - H, I, J, K, L, M, N, O are terminal states
 - Start: A, Goal: O











COMPLEXITY

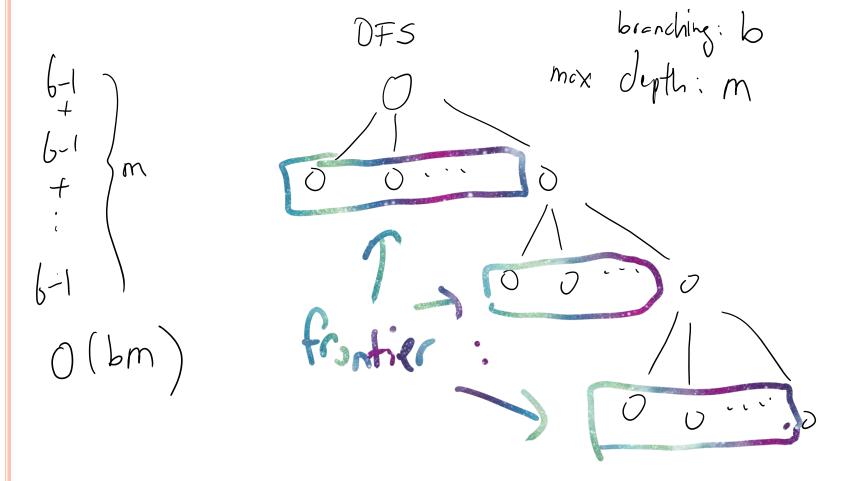
- We're concerned with
 - Time complexity
 - Space complexity
 - Completeness (if there is a solution, can it find it?)
 - Optimality (is the found solution the optimal one?)
- Three relevant quantities
 - The number of nodes expanded
 - The maximum size of the **frontier**
 - The cost of the solution

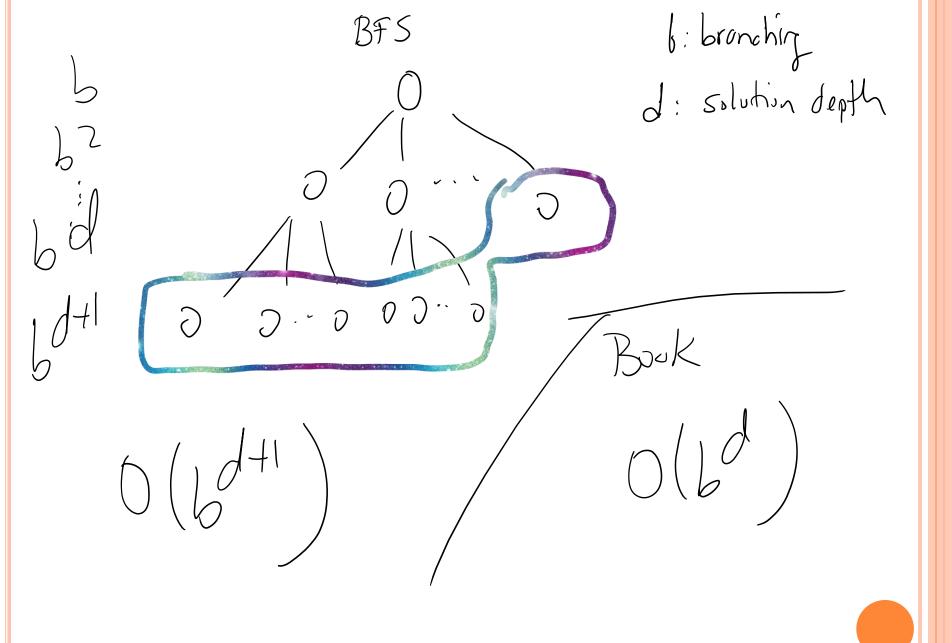
Analysis – Tree Search

Criterion	Depth-first	Breadth-first
Complete?		
Time		
Space		
Optimal?		

Analysis – tree Search

Criterion	Depth-first	Breadth-first
Complete?	No	765
Time	$O(P_{W41})$	0(6)
Space	O(pw)	$5(b^{d+1})$
Optimal?	No	Yes, if all step
		have equel cost



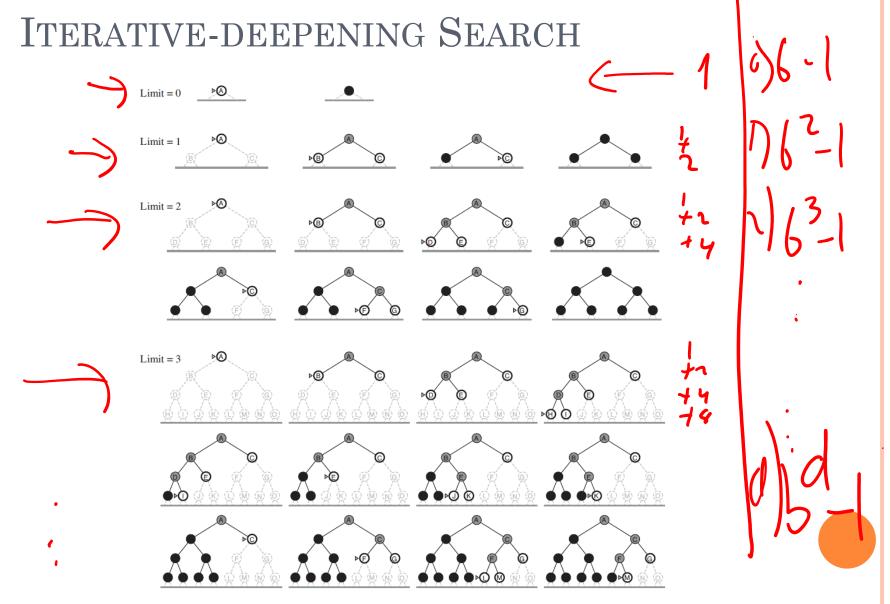


MEMORY IS A BIG PROBLEM

- Breadth-first search has an exponential memory requirement
- Depth-first search requires polynomial size memory,
 but it is not complete and it is not optimal
- What can we do?
 - Depth-limited search
 - o Impose a limit on the depth
 - o Is it complete? Optimal?
 - Iterative-deepening search
 - Start with a depth limit of 0, and run depth-limited search; if no goal found, increase the depth limit and repeat the process
 - o Is it complete? Optimal?

This is the textbook version; we do the DFS part differently.

What are the time and space complexities for IDS?



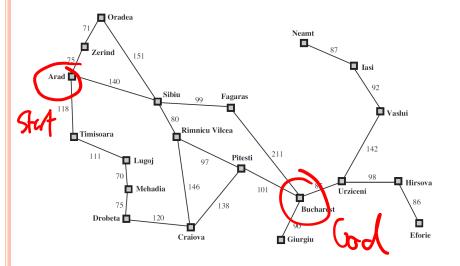
WHAT IF THE STEP COSTS ARE NOT IDENTICAL?

- BFS and Iterative Deepening are not guaranteed to be optimal anymore
- Define f(n) = h(n) + g(n)
 - h(n): Cost estimate from n to the goal
 - g(n): Cost from the initial state to n
- 1. Uniform-cost search
 - Expand using g(n)
- 2. Greedy best-first search
 - Expand using h(n)
- 3. A* search
 - Expand using f(n)

THE RUNNING EXAMPLE

Edge costs

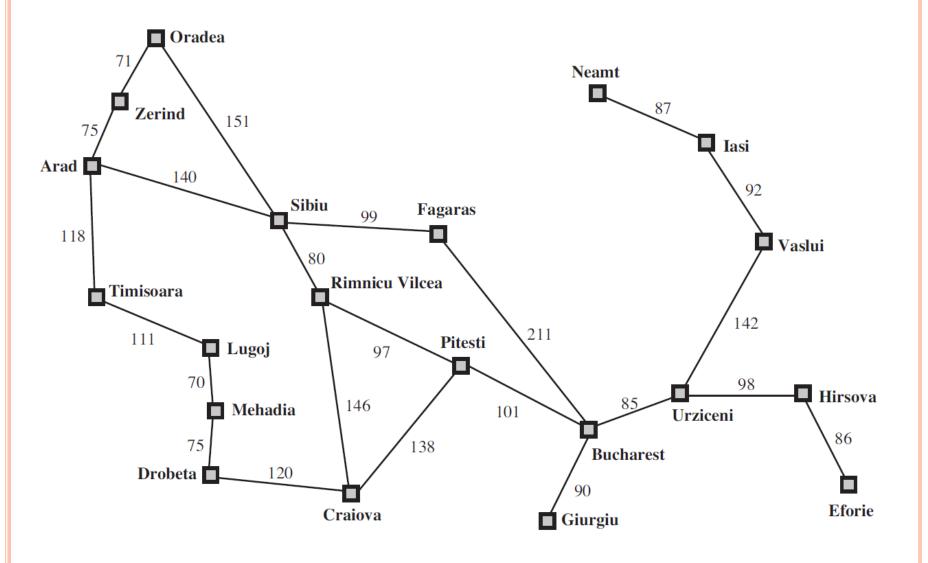
h(n)



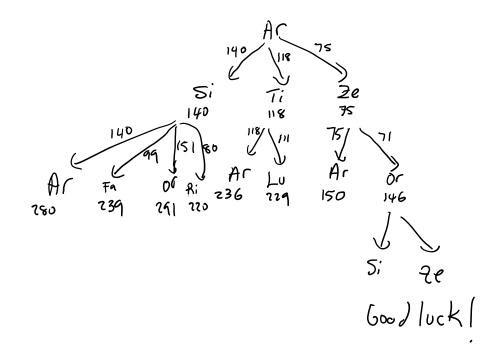
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193 —
Fagaras	176—	Sibiu	253
Giurgiu	77	Timisoara	329 /
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374 —

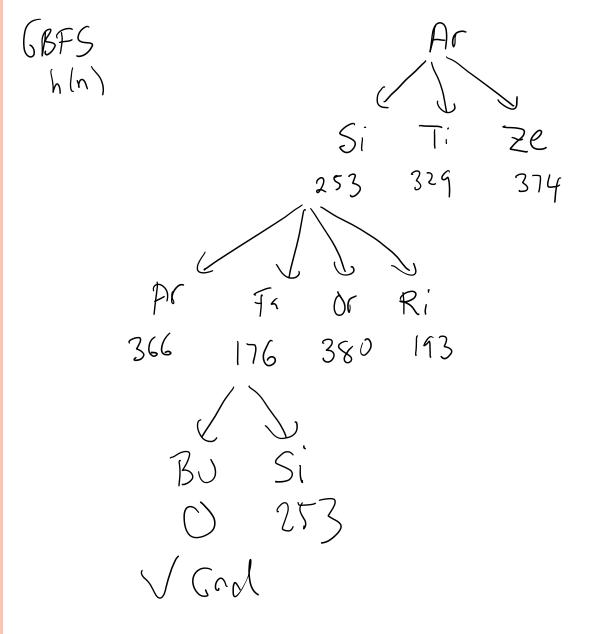
UNIFORM-COST SEARCH

• Arad to Bucharest

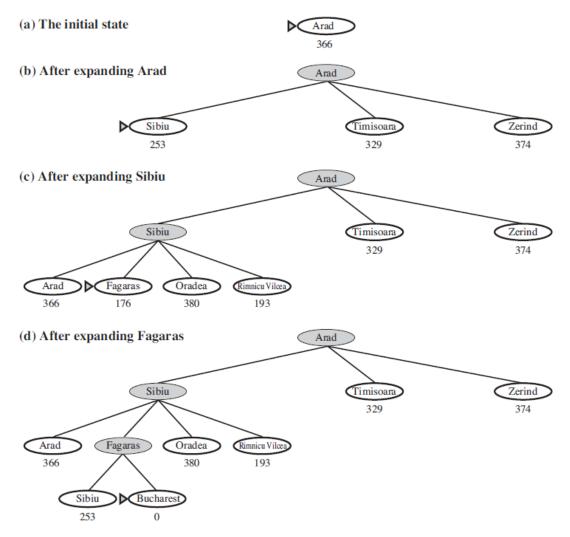


UCS g(n)

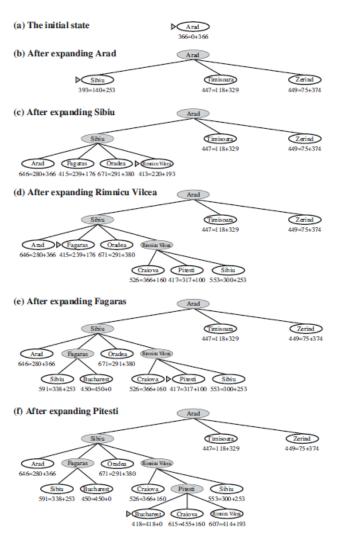




GREEDY BEST-FIRST SEARCH



A* SEARCH



GOAL TEST

Test when a goal state is generated or chosen for expansion?

REPEATED STATES

- Repeated states can be a major problem
- Repeated states are especially common for problems where actions are reversible
 - But, they also occur for problems where there are multiple distinct paths between two states
- One way to avoid this problem is to remember the set of explored and generated states

TREE SEARCH VS. GRAPH SEARCH

function TREE-SEARCH(problem) returns a solution, or failure
initialize the frontier using the initial state of problem
loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier if the node contains a goal state then return the corresponding solution expand the chosen node, adding the resulting nodes to the frontier

function GRAPH-SEARCH(problem) **returns** a solution, or failure initialize the frontier using the initial state of problem

initialize the explored set to be empty

loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier

if the node contains a goal state then return the corresponding solution

add the node to the explored set

expand the chosen node, adding the resulting nodes to the frontier only if not in the frontier or explored set



Modified Graph Search

function GRAPH-SEARCH(problem) returns a solution, or failure initialize the frontier using the initial state of problem initialize the explored set to be empty loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier if the node contains a goal state then return the corresponding solution add the node to the explored set expand the chosen node

for child in children of the chosen node:

if child is not in explored:

add child to frontier only if child is not in frontier or child has a lower cost than the one in frontier

GRAPH SEARCH EXAMPLES

QUESTIONS TO THINK ABOUT

How does graph search affect

- Completeness,
- Time complexity,
- Space complexity, and
- Optimality

of various search algorithms?

A* OPTIMAL?

- Tree version is optimal if
 - h(n) is admissible
- Graph version is optimal if
 - h(n) is consistent

ADMISSIBLE? CONSISTENT?

- Admissible if
 - h(n) never overestimates the optimal cost
- Consistent if
 - If $h(n) \le c(n,n') + h(n')$

A* Optimality – Graph Search – Proof

- 1. If h(n) is consistent, then f(n) along any path is non-decreasing
- 2. When *n* is expanded, the optimal path to it has been found
- Proof is given on page 95

A* VS UNIFORM-COST SEARCH

- Compare and contrast A* and UCS
 - Which one is better and why?
- Given the same heuristic function, can you design a complete and optimal algorithm that expands fewer nodes than A*?

8-PUZZLE HEURISTICS

- \circ h_1 = The number of misplaced tiles
- $o h_2 = Manhattan distance$

Let's see an example

5	6	1	
2	4	8	
3	7		

EFFECTIVE BRANCHING FACTOR

- If the total number of nodes generated is N and the solution depth is d, then
 - b* is the branching factor that a uniform tree of depth d would need to have in order to contain N+1 nodes
- $N+1 = 1 + b^* + (b^*)^2 + ... + (b^*)^d$
- o If A^* finds a solution at depth 4 using 40 nodes, what is b^* ?
 - ≈2.182
- A good heuristic function achieves $b^* \approx 1$

IDS vs A*

	Search Cost (nodes generated)		Effective Branching Factor			
d	IDS	$A^*(h_1)$	A*(h2)	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14		539	113		1.44	1.23
16		1301	211	r ha ryhania	1.45	1.25
18		3056	363	to trotte of	1.46	1.26
20		7276	676		1.47	1.27
22		18094	1219		1.48	1.28
24		39135	1641	cell in His	1.48	1.26

Figure 3.29 Comparison of the search costs and effective branching factors for the ITERATIVE-DEEPENING-SEARCH and A* algorithms with h_1 , h_2 . Data are averaged over 100 instances of the 8-puzzle for each of various solution lengths d.

HEURISTIC FUNCTIONS

- Let h(n)=0 for all n
 - Is it admissible?
 - Is it consistent?
 - Is it any good?
- o Can we say an admissible heuristic function h_i is always better than another admissible heuristic function h_i ?
- o How can we find good heuristics?
- A heuristic: run uniform-cost search, find the solution, return the cost of the solution as the heuristic value. Can you beat it?
- Let $h(n) = \max(h_1(n), h_2(n), ..., h_m(n))$, where $h_i(n)$, are admissible
 - Is h(n) admissible?
 - Can we say h(n) is better than any of $h_i(n)$?
 - What's the catch?

SUMMARY

- Depth-first search
- Breadth-first search
- Depth-limited depth-first search
- Iterative deepening
- Uniform cost search: g(n)
- \circ Greedy best-first search: h(n)
- A* search: g(n) + h(n)

Which one is better under what conditions?

NEXT

- Skip Chapter 4
- Chapter 5: Adversarial search game playing